NoHandsHam — release notes

Version 1.3.2

Speech to text for amateur radio operators: it listens to your radio, writes down what was said, and pulls the call signs out of it. Windows, and it runs on the machine rather than needing a browser open.

Download the newest installer

This repository holds no source code. It exists to publish the installer and the manuals, so that downloading NoHandsHam does not mean going anywhere near a source tree. The program itself lives in michael-h-peters/NoHandsHam.

Below is every release and what was in it, newest first. Each entry is a line or two; the full notes for a release are on its own page, and the manual explains any of it properly. There is a PDF of this page if you would rather keep one.

Releases from 0.13.2 are downloadable here. Everything before that predates this repository and is listed for the record — the version numbers are what the program reports at the foot of its window, so an old copy can still be placed.

What is in a release

file	what it is
`NoHandsHam-<version>-x64.exe`	the installer — this is the one to download
`GETTING-STARTED.pdf`	the short guide: install it, get a key, make a first entry
`README.pdf`	the full manual, every setting and every flag

Both PDFs are installed alongside the program as well, and reachable from the **Documents** button in the window, so there is nothing to keep track of after the install.

Double-click the `.exe` and answer the wizard. It carries everything it needs, including — from 1.1.8 — Microsoft's Visual C++ runtime, which the on-machine transcriber depends on and which is skipped in a second if you already have it. For an install with no clicking:

```
.\NoHandsHam-1.3.2-x64.exe /quiet                    # silent
.\NoHandsHam-1.3.2-x64.exe /quiet ADDTOPATH=1        # also put nohandsham.exe on the PATH
.\NoHandsHam-1.3.2-x64.exe /quiet DESKTOPSHORTCUT=0  # skip the desktop shortcut
.\NoHandsHam-1.3.2-x64.exe /uninstall /quiet         # remove it
```

Upgrades replace the installed copy where it stands. Nothing to uninstall first, and your settings, keys, licence and transcripts live under `%AppData%` and are not touched by an install, an upgrade or a removal.

The version is at the foot of the window beside the copyright, and is the first thing worth quoting in a bug report. The program checks once a day whether a newer release has been published here and says one line when there is one. It never downloads or installs anything on your behalf, and unticking **Check version** on the top strip of the window switches the checking off for good.

Releases

1.3.2 — 30 October 2026

Fixed: phonetic call sign detection. Newer AI models often return transcriptions with hyphenated phonetic words (e.g., `whiskey-one-alpha-whiskey`), which NoHandsHam now recognizes.

- **Fixed:** Hyphenated runs of phonetic words and digits are now correctly parsed into call signs.
- **Fixed:** Hyphenated multi-word tokens like `double-u` are correctly identified as symbols.

1.3.0 — 29 October 2026

Removed: Vosk, and the local draft check. NoHandsHam transcribes in two places again — at OpenAI, or with whisper.cpp on this computer. That is the choice it offered through 1.1.9, and the one it offers now.

- **Removed:** *This computer — free and quick*, the Vosk engine added in 1.2.0. The **Transcription** row has two options again, and `-vosk` and `-vosk-model` are no longer flags.
- **Removed: Check each utterance here first**, the Moonshine pre-pass added in 1.1.10 and given a settings row of its own in 1.2.1. The **Draft check** row goes with it, along with `-draft`, `-draft-model`, `-no-draft-gate` and `-no-draft-hint`.

If you were using Vosk, this moves you to whisper.cpp rather than to OpenAI. Your settings file names an engine this version does not have, and the ordinary rule for that — fall back to OpenAI — is the one wrong answer here, because somebody who chose Vosk chose it to avoid needing an API key. whisper.cpp answers the same question: free, offline, no key, nothing uploaded. You will need to press **Download** once, and the line above the button says the size first. The diagnostics log records that the engine was changed.

If you were using the draft check, nothing replaces it. What it was best at — stopping a burst of static being sent, billed, and coming back as a call sign nobody said — is the **Gate** and **Margin** sliders' job alone again. If invented call signs were what it was saving you from, **Call signs to ignore** is the other tool for that, and raising the **Gate** toward −45 is the first thing to try.

Your downloads are left alone. Vosk and Moonshine stay on disk in `%AppData%\whisper-listener\vosk` and `...\moonshine` — about 500 MB between them if you had both. Nothing uses them now, and

NoHandsHam does not delete files you downloaded; both folders are safe to remove by hand whenever you like. Your whisper.cpp install, in the `whisper` folder beside them, is untouched.

Everything else is as it was. No setting has changed its meaning or its default, the settings pane is shorter by the two rows that went, and nothing about how utterances are heard, cut up, transcribed or logged has moved.

1.2.1 — 29 October 2026

Changed: the draft check is its own setting. Nothing about how NoHandsHam works has moved — this is about being able to tell two things apart on the settings pane.

The **Transcription** row chooses where your words are worked out: OpenAI, whisper.cpp or Vosk. **Check each utterance here first** is not one of those — it is Moonshine, a small model that listens to each utterance before it is sent to decide whether it is worth sending at all, and it runs in front of whichever engine you picked. The two have always been separate installs with separate downloads, separate folders and separate Download buttons, and choosing an engine has never touched the check or the other way round. The pane did not say so: with *This computer — free and offline* selected and the check ticked, one row carried two model pulldowns, two Download buttons, two progress bars and two status lines, told apart only by a caption reading **Draft:** where the other read **Model:**.

- **Changed:** the check and everything belonging to it — the two tick boxes, the model pulldown, its Download button and its status line — now sit on a row of their own called **Draft check**, below **Base URL**. The engine keeps the **Transcription** row.
- **Changed:** the status line names **Moonshine**, so the two Download buttons can be told apart by reading rather than by guessing. Before the download it names sherpa-onnx as well, the toolkit that runs the model — on its own beside a name like `tiny.en` it read as an unexplained third thing being fetched.
- **Changed:** both status lines say in words that this is a separate install from the engine above, and the **?** beside each row explains the difference from its own side.

Draft check sits below **Model** and **Base URL** rather than directly under the radio buttons because those two boxes belong to the engine — it takes both over and greys them — so they finish the thought the **Transcription** row starts. A row slipped between a question and its own answer would read worse than the arrangement it replaced.

The settings pane is the same width it was and the same height as it comes up, so nothing has been pushed below the fold to pay for this. Your settings carry over untouched: if you had the check on with a model chosen, it is still on, with the same model, in the same place on disk. No setting has changed its meaning or its default, and nothing about transcription itself has changed.

1.2.0 — 14 September 2026

New: a third place to transcribe. NoHandsHam could already work the words out at OpenAI or on your own computer with whisper.cpp. It can now do it with **Vosk**, which is the light end of the range — about 53 MB to fetch against whisper.cpp's 473, and quick enough that the transcript arrives while the next station is still speaking.

- **New:** *This computer — free and quick* on the **Transcription** row. Nothing is uploaded, nothing is billed, no API key is needed and a net can be logged with the cable out — the same bargain as whisper.cpp, at a tenth of the download. On the machine this was built on it transcribed seven and a bit seconds of speech in under a second, roughly an eighth of the time the speech itself took. No graphics card and no Microsoft runtime needed.
- **New:** two models, both English-only. **small-en-us-0.15** (39 MB) is the one to use and is the whole point of this engine; **en-us-0.22-lgraph** (125 MB) makes fewer mistakes for three times the download.
- **New:** `-vosk` does the same from a script, with `-vosk-model`. Asking for `-local` and `-vosk` in one run is an error rather than a guess at which you meant.
- **New:** the top strip says `Engine: Vosk` while it is selected, beside the `Engine: OpenAI`, `Engine: CPU` and `Engine: GPU` it already said.
- **Fixed:** the `Model:` caption beside the whisper.cpp model pulldown was shown even with OpenAI selected, standing on its own with nothing beside it. It now goes away with the row it belongs to.

What it costs is worth knowing before you switch. Vosk mis-hears more than whisper.cpp does, and for this program that means call signs — the error that matters most and the one you are least likely to notice. It writes in lower case with no punctuation at all, so a transcript reads like a telegram; call signs are still picked out and logged properly, because NoHandsHam reads those out of the phonetic alphabet and that does not need capitals. Your roster of known call signs does not help here the way it helps OpenAI, because Vosk cannot be told what it is about to hear — **Call signs to ignore** still works, and matters more.

It is also the one engine that runs inside NoHandsHam rather than as a separate program. There is no port, nothing listening on your network and nothing left running if the app closes; the other side of it is that if Vosk itself falls over it takes NoHandsHam with it, where a crash in the whisper.cpp server costs one utterance.

Download is the consent, as it is for the other engines — nothing arrives because you pressed a radio button, and the line under the pulldown says the size first. It lands in `%AppData%\whisper-listener\vosk`, which an upgrade leaves alone, and the two local engines share nothing: trying Vosk for an evening leaves your whisper.cpp installation exactly where it was. Vosk is Alpha Cephei's, under the Apache 2.0 licence, and so are both models; NoHandsHam includes none of it. Nothing else has changed, no existing setting has changed its meaning or its default, and OpenAI is still what a fresh install uses.

1.1.10 — 13 September 2026

New: NoHandsHam can listen to each utterance on your own computer first, and drop the ones that turned out to be noise rather than speech. Off until you tick it.

- **New: Check each utterance here first**, at the foot of the *Transcription* row. NoHandsHam decides somebody is talking by how loud the audio is, because that is the only thing there is to measure before the words exist — so a door closing, a fan, or a burst of static on a quiet band can be loud enough to count, and then that noise is sent off to be transcribed and comes back as words nobody said, very often as a call sign. With this on, every utterance is first transcribed here by a small local model, and what it heard is used two ways, each with its own tick box: utterances it found no words in are **not sent at all** — not transcribed, not billed, not written down — and the rest are sent with the draft as a prompt, so the real transcriber starts with a rough idea of what it is about to hear. Fed three seconds each of digital silence, mains hum and white noise, the local model returns nothing at all for all three, which is what makes the first half work.
- **New:** `-draft` does the same from a script, with `-draft-model`, `-no-draft-gate` and `-no-draft-hint`.
- **New:** if the check ever drops everything, the session line says so outright rather than leaving you with an open microphone and an empty window — *nothing transcribed: the local check found no speech in 8 utterances · untick it in Settings if that is wrong*. The CLI prints the count when it stops.

What it costs is about three-quarters of a second of your processor per utterance, before the utterance is sent, and nearly all of that is loading the model rather than listening to you. 147 MB the first time, and **Download** is the consent — nothing arrives because you ticked a box. Two models are offered, both English-only: **tiny.en** (124 MB) is the one to use, since it only has to spot the words and never to spell them, and **base.en** (287 MB) drafts better in twice the time. If you transcribe at OpenAI the first half pays for itself, because every utterance of static it stops is a minute you are not billed for; if you transcribe on this computer nothing is billed either way and what it buys you is a log without invented call signs in it.

The models are Moonshine, under the MIT licence, run by sherpa-onnx under Apache 2.0. NoHandsHam includes neither — pinned releases of both are fetched on the first Download, checked against a recorded fingerprint, and run unchanged. Nothing listens on your network and nothing is left running: the check is one short-lived program per utterance, gone when it has looked. Nothing else has changed, and no existing setting has changed its meaning or its default.

1.1.9 — 27 August 2026

If the right-hand end of the settings pane looked cut off, this is the fix. The band box showing "1.25" of "1.25m" is the one most people saw.

- **Fixed:** the settings pane is now held clear of its own scrollbar. The pane is always taller than the window, so it always has a bar down its right-hand side, and that bar is painted on top of the pane rather than beside it. Every row with a control at its right-hand end sat flush against that edge, so the bar covered part of it — twelve units while the mouse was over the pane, six while it was not. Both **Rescan** buttons, **Browse…**, **Default** and the **Band** box were affected. The band box is much the narrowest at fifty-seven units, so it lost about a fifth of itself, which is why it was the one that got reported. One inset in the place the whole pane passes through holds every row clear at once; nothing in the app had allowed for the bar anywhere before.

This was reported as the right side being chopped off on a 1920x1080 laptop when the window was made smaller, and the window size turned out to have nothing to do with it: the loss was the same at every width and on every screen. Making the window narrower squeezes the frequency box beside it, which crowds that end of the row and is what draws the eye. Nothing else in the settings has moved and no setting has changed its meaning or its default.

1.1.8 — 27 August 2026

If pressing *Start listening* told you "the transcription server stopped before it was ready", this is the fix. It only ever happened when transcribing *on this computer*; the OpenAI engine was never affected.

- **Fixed:** the installer now carries Microsoft's Visual C++ runtime. whisper.cpp, the transcriber that runs on your own machine, is built with Microsoft's compiler and needs that runtime; Windows does not include it, the whisper.cpp download does not contain it, and until now neither did this installer. On a computer that had never installed another program built the same way, Windows killed the transcription server before a line of its own code ran — so it had no chance to say anything, which is why the message was one sentence and the log pane under it was empty. Any machine that has installed a game, Office or almost any other desktop program already had the runtime and never saw this.
- **New:** if the runtime is missing, the *Transcription* row says so and names where to get it — **before** the download, so a machine that cannot run the transcriber does not spend 671 MB finding that out. It is never removed when you uninstall NoHandsHam, because other programs depend on the same runtime.
- **Changed:** if the server fails to start for any reason, the message now carries the exit status Windows gave and the last thing the server managed to print. The two statuses that mean a missing runtime are named in plain words.
- **Fixed:** settings you had changed could be replaced with the defaults. NoHandsHam could not tell "no settings file yet" from "a settings file I failed to read", so on a run where reading went wrong it treated the machine as brand new, and the ten-second autosave then wrote defaults over the file it had just failed to read. A settings file that cannot be read is never overwritten now.
- **The installer is about 49 MB rather than 23.** Microsoft's redistributable is half of it, it arrives already compressed, and carrying it is the difference between the transcriber working on a clean machine and not.

Already on 1.1.7 and seeing this? Installing 1.1.8 fixes it. If you would rather not reinstall, the runtime is a 25 MB download from Microsoft at aka.ms/vs/17/release/vc_redist.x64.exe — install it, reopen NoHandsHam, and the server starts.

1.1.7 — 27 August 2026

The four sliders start somewhere different. If you have never touched them this changes what NoHandsHam does; if you have, your settings are kept and none of it reaches you.

slider	was	now
Gain	0	0 — unchanged
Utterance length	25 s	**10 s**
Gate	−42 dBFS	**−60 dBFS**
Margin above room noise	+8 dB	**0 dB**

- **Changed:** the gate is open. The two sensitivity sliders together mean speech has to beat whichever is higher of −60 dBFS and the measured room noise — and with the margin at 0, that is the room noise itself with no gap above it. The old +8 dB gap was protecting against something a radio feed does not have: this program is for a receiver patched in on a cable, a line feed has no room noise, and the gap only cost you the quiet stations. It does not switch the adapting off — the limit still follows the noise up and down as the band changes, it just sits level with it.
- **Changed:** an utterance is cut after 10 seconds rather than 25. A turn is transcribed only once it ends, so that ceiling is the longest you can wait to see a line of text; a talker who never pauses now appears in ten seconds. The price is the model seeing less of the turn at once, which is why it stays a slider.
- **The case these are wrong for is an open microphone in front of a loudspeaker** — the gate will trip on the room. Raise **Margin** to +4, then +8. Both manuals now say so where you would look for it.
- **Fixed:** fifteen places in the manuals quoted the old numbers, and some advice had become nonsense against the new floor — "*missing weak stations? lower the Gate toward −50*" when it now starts at −60. Both slider walkthroughs, both worked examples of the live reading, the chunking formula, the too-hot-and-too-cold section and both symptom tables are rewritten rather than renumbered.
- **Fixed:** two sentences still said settings are written only when you close the window, which stopped being true when 1.1.4 started saving them every ten seconds.

1.1.6 — 27 August 2026

If the right-hand side of the window has been off the edge of your laptop screen, this is the fix.

- **Fixed:** the window opens at a size the screen can actually show. It was not about the width of anything inside it — the window was opening bigger than the display. NoHandsHam measures the screen at start-up and cuts its opening size down to fit, but it decided both dimensions together and abandoned the cut on *both* if either looked implausible. An ordinary laptop is too short by that measure: a 1366×768 panel at 150% scaling has about 450 of the units the app counts in, against a threshold of 560. So the measurement was rejected, the width cut went with it, and the window opened at its full 880 units on a screen 781 wide. Only 100% scaling escaped. Each dimension is judged on its own now, against what a real desktop reports rather than what a roomy window would prefer.
- Seven ordinary laptop configurations were checked, from 1366×768 at 100% through 1920×1080 at 250%. All of them now open inside the screen, title bar and borders included.
- Why it took a few goes: 1.1.4 and 1.1.5 made the window and the settings inside it able to *become* small, and both were real fixes — the window can be dragged to 454 units and the form fits in 436,

which is what makes this cut safe at any screen size. Neither addressed the window still *opening* too large, which is what put the right edge off the display.

1.1.5 — 27 August 2026

- **Fixed:** the settings now fit inside the window, however small you make it. 1.1.4 let the window be dragged narrow but left the form inside it wanting 719 units of width, so on a laptop the settings had to be scrolled sideways to be read. The form asks for 436 now — less than the narrowest the window can be — so the sideways scrollbar is there and never needed, and no setting is cut off at any size.
- **Changed:** the station row is two rows — *Call sign* and *Location* on the first, *Freq* and *Band* on the second. *Park/Location* is now **Location** and *Frequency* is now **Freq**, and each box is sized to what goes in it: a POTA reference is `K-1234`, a band is `1.25m`. The call sign box keeps room for a suffixed call like `KD2ABC/P`.
- **Changed:** three tick boxes had whole sentences for labels. They have short labels now with the sentence on the line underneath — moved, not dropped. *Send a radio word list* keeps its guidance about stray words that way, and *Police phonetics* keeps "Adam Boy Charles", which is the part you recognise. The two transcription options lost only words the help line directly beneath them already said. Nothing was removed and no setting moved section.
- **Fixed:** two faults in the wrapping rows 1.1.4 introduced, which were costing 78 units of blank space below the top of the settings — a height that could only grow once it had been measured at a narrow width, and a wrapped row that never told the layout above it that it had grown.

1.1.4 — 27 August 2026

Two bug fixes, both worth taking.

- **Fixed:** settings no longer disappear when you upgrade. Most of what you set — the call sign, the mode, the microphone, every slider, the transcription engine, and the window's own size and position — used to be written only when you closed the window. An upgrade never closes it politely: Windows Installer stops the running copy so it can replace the files underneath it, and everything changed since launch went with it. So did a crash's worth, or a flat battery's. Settings are now written while you work: the pane is checked against the file every ten seconds and saved when the two differ, so the most any of them can cost is the last ten seconds. An untouched window writes nothing at all. The API key and the QRZ password were never affected — they are sealed separately and were always saved on the press.
- **Fixed:** the window can be made small. A window cannot be dragged narrower than its contents need, and the widest row that could not break held it at 741 units — wider than a 13-inch laptop has to spare at 200% display scaling, so on those machines it opened wider than the screen and could not be shrunk at all. The top strip and the row of buttons at the foot now wrap onto a second line instead of holding the window open, and the settings pane scrolls sideways rather than setting a floor of its own. Where 1.1.3 snapped back to 794 pixels whatever you asked for, 1.1.4 takes 520.

- The trade on that second one: pull the window in far enough and it needs to be a little taller, because those rows have wrapped and genuinely need the line. Nothing moves at ordinary sizes.
- **Changed:** the engine badge sits after the **?** on the top strip rather than in the middle of it, which is what lets that row wrap.

1.1.3 — 26 August 2026

If you are on 1.1.1 or 1.1.2, take this one.

- **Fixed:** the main window resized itself and then refused to be resized. The engine badge 1.1.1 added to the top strip is a text label, and a text label asks for exactly the width its current words need — which the toolkit enforces on the window in both directions, growing one that is too narrow and refusing to let you drag below it. So the narrowest the window could be went from 719px to 836px, and a window left narrower than that was pushed out at launch and would not go back.
- **Fixed:** worse, the limit *moved*. The three captions are different lengths, so it was 836px reading `Engine: OpenAI` and 814px reading `Engine: CPU` — switching engine resized the window, and a width you could drag to a minute ago you suddenly could not. The badge is also rewritten when you tick *Use the graphics card* and whenever the transcription server starts or stops, none of which you would connect to the top strip, which is why it looked random.
- The badge may now shorten with an ellipsis, which uncouples its width from its text. The floor is 741px and no longer moves — within 22px of what it was before any of this existed, checked across every engine and graphics-card state. The caption still shows in full unless you drag the window right in.
- Nothing else changes: not how transcription works, not what gets downloaded, not what the badge says.

1.1.2 — 25 August 2026

- **Fixed:** the **Use the graphics card** tick-box explains itself instead of disappearing. On a computer with no NVIDIA card — most computers, and every machine with AMD or Intel graphics — 1.1.1 simply left it out, and nothing anywhere else mentioned graphics either, so the feature and the absence of the feature looked identical. A control that is missing cannot say why it is missing. It is shown greyed now, reading *no NVIDIA card on this computer* or *its driver is too old* in its own label, which is where anyone hunting for it is looking.
- **Fixed:** a box that cannot be ticked no longer shows a tick. Settings move between machines, and a ticked box on a computer with no card read as "the card is in use" on hardware that has none.
- **Fixed:** a choice you were never offered is no longer recorded as one. Whether the preference is saved was decided by asking whether the box was on screen, and now that it is always on screen every machine without a card would have written "not the graphics card" — which, restored onto a desktop that has one, would have arrived having already declined it.
- Nothing changes about how transcription works on either engine, or about what gets downloaded.

1.1.1 — 25 August 2026

- **New:** the top strip of the window says which engine is in use — `Engine: OpenAI`, `Engine: CPU` or `Engine: GPU`. 1.1.0 added the graphics card and then said so only in the *Transcription* row, which is on a settings pane taller than the window: "is my card actually being used" cost a scroll and a read, which is a poor answer to a question you ask precisely because you are not sure. It is now visible with *Settings* collapsed and needs nothing pressed.
- It reports the **running server**, not the tick-box. The transcription server stays up between sessions so that pressing *Start listening* again does not reload half a gigabyte, so ticking *Use the graphics card* mid-session leaves the badge reading `Engine: CPU` until the process it describes really is the card's. A badge that followed the tick-box would have been claiming hardware that was not doing the work.
- `CPU` and `GPU` rather than the *processor* and *graphics card* the pane spells out: everything on that strip competes with how narrow the window can be dragged, and this is the first thing up there to set that minimum — 719px before, 836px now, against the 880px the window opens at. The longer phrasing would have left 3px of that margin.
- Nothing changes about how transcription works on either engine. This release only says what was already true.

1.1.0 — 25 August 2026

- **New: transcribe on the graphics card.** Transcribing on this computer arrived in 1.0.0 and has run on the processor ever since. If the machine has an NVIDIA card, a **Use the graphics card** box appears in the *Transcription* row and the work goes there instead — quicker at the arithmetic that is most of the job, and it leaves the processor free for the microphone, which matters because audio capture runs in the same program and a dropped buffer loses words for good. It is the same model either way: faster, not more accurate.
- **What it costs, named before you press anything.** There is one ready-made Windows build of whisper.cpp for NVIDIA and it carries NVIDIA's own maths libraries inside it, so nothing has to be installed separately: 671 MB to download against the processor build's 8 MB, and about a gigabyte on disk. The line under the model pulldown names *both* figures first, and the tick-box is the consent — the same rule the model download has always followed.
- **An installation that already works is left alone.** Upgrading from 1.0.0 does not replace a working processor build with a 671 MB download. The card is mentioned, with what it would cost, and the choice stays yours. The two builds sit in separate folders and share the speech models, so ticking the box never re-downloads a model.
- **NVIDIA only, and only on Windows**, because whisper.cpp publishes no ready-made Windows build for AMD or Intel graphics — a limit of what there is to download, not a judgement about the hardware. On those machines the box does not appear. The driver must be 528.33 or newer; an older one is named alongside the version needed, since that is usually a ten-minute update rather than a dead end.
- **New:** at the command line, `-local-cpu` declines the card, for a metered connection or a small disk. `-local` on its own picks it up when there is one.

- **Changed:** the **Running** line on the settings pane describes the server that is actually running rather than whatever the pulldowns currently say. The server stays up between sessions on purpose, so that pressing *Start listening* again does not reload half a gigabyte — which meant changing the model mid-evening left the pane describing one thing and a process doing another. It now names the model that is loaded and says when your choice has moved on from it.

1.0.0 — 23 August 2026

- **New: transcribe on this computer**, for nothing a minute. Every minute of audio used to be billed — by the minute of *audio*, not by the word, so a three-hour net that was mostly silence cost three hours. A new **Transcription** row on the settings pane offers *This computer* instead of *OpenAI*, and with it chosen no API key is used, nothing is sent anywhere, and a session works with the internet unplugged.
- **New:** the server and the model are fetched for you, once, when you press **Download** — not because you clicked a radio button, because half a gigabyte is not a side effect. Four models to choose from: tiny.en 74 MB, base.en 141 MB, small.en 465 MB (the default, and the best trade for radio) and medium.en 1.4 GB. They live under `%AppData%`, which an install never touches, so an upgrade does not re-download them.
- **What it costs**, said in the pane and in the manual rather than found out during a net: call signs through noise are harder for a model on your desk than for the hosted one, and there is no live partial text — text arrives one finished utterance at a time. The status line does not claim otherwise.
- The server is whisper.cpp, MIT-licensed and fetched from upstream, pinned by version and checked against a SHA-256 recorded in the program — this downloads an executable and then runs it. It binds to loopback so nothing on your network can reach it, opens no console window, and closes when NoHandsHam does.
- **Changed:** with *This computer* chosen the API key line reads **not in use** rather than reporting a fault. The local server has no key to check, and a working setup should not be labelled rejected.
- **New:** at the command line, `-local` and `-local-model`. The window and the command line have always been the same program underneath, and neither should be able to do something the other cannot.
- OpenAI stays the default, and nothing about an existing install changes unless you touch the new row.

0.20.0 — 23 August 2026

- **New:** checking a station into the net log is one click. Every entry in the transcript that names a station now carries a second box beside the roster's, and ticking it is the whole check-in — at the time the entry was heard, with the number the station took said at the foot of the window. It used to cost a dialog: click the entry, wait for the editor, tick *In the net log*, press **Save**.
- **New:** one station is usually on several entries, and they all agree. Tick the box on any of them and every other entry naming that station shows ticked; untick any one and they all clear. The net log window need not be open for it, or at all — tick your way through the net and read the list at the end.

- *In the net log* is still in the entry editor, for the check-in you make while putting a mis-heard call sign right. Both boxes are the same act.
- **New: call signs to ignore**, on the Mode row and kept between runs — for the repeater identifier, the club name read out, the frequency spoken as letters. One on that list is dropped where it is first heard rather than tidied up afterwards: it never reaches the transcript, the roster, the prompt, a QRZ lookup or a check-in box.
- **Changed:** the version check moved to the top strip of the control window, beside *Getting started* and *Manual*. It was the last row of the settings pane, which is taller than the window, so on a laptop-fitted window it sat the better part of eight hundred pixels below the fold — the one launch in a month that had news put the news where nobody was looking.

0.19.0 — 23 August 2026

- **Fixed:** running a second net without pressing **Clear** first published it under the *first* net's identifier, writing over it. Three nets went up as one row. Clear now always ends the net, including on a list already emptied by hand, which was the hole behind it.
- **New:** the publish line names the net and when it started — *Published "Tuesday Traffic Net" to nohands.net at 20:14:07, net started 19:58*. The start time is the tell: a net that says it began this morning when you called this one ten minutes ago is the fault above, visible without opening a browser.
- **Changed:** everyone signed in to nohands.net can now read every net published to it, including the stations who checked in. Accounts there are approved by hand, so it is a shared log among approved accounts on one site rather than the open web; there is still no public link, and deleting a net you published is yours alone.
- **Changed:** because that replaces a promise the old wording made, the dialog goes up once more with the new text if you already have publishing switched on. Keep publishing or stop; stopping leaves what is already on the server where it is.

0.18.0 — 23 August 2026

- **New: Publish to nohands.net** — the net log can go on the web while the net runs, updating as the list changes. Sign in at nohands.net to see it; only you can, and there is no public link.
- **Off unless you tick it**, and it asks first. This is the first thing NoHandsHam sends anywhere that is not yours — other stations' call signs, names and towns — so upgrading never switches it on.
- Needs a licence key; a copy on trial cannot publish.
- **New:** a dropped connection costs nothing. The net carries on, the window says it is still trying, and the next send catches the server up.

0.17.1 — 23 August 2026

- **Changed:** the net log's small print is behind a **Help** pull-down at the top of the window — three topics, plus the manual. It was a paragraph above the table, read once and then charging four lines of table for

the life of the install.
- **Changed:** the topics say *more* than the paragraph did. A dialog has room for what a strip above a table has to cut.
- The **Save…** reminder stayed on the window, beside the count. Of the four things that paragraph said it is the only one with a deadline, and learning it from a menu you never opened means learning it after closing the app on an evening's work.
- About twelve check-ins are visible at once, up from eleven.

0.17.0 — 23 August 2026

- **New:** the net log has a header — the net's name, net control's call sign, name and QTH, and the frequency and band. All of it is written into the saved file.
- **New:** the call sign, frequency and band in that header *are* the ones on the settings pane. Change either and both follow, so you can retune from whichever window is in front of you.
- **New:** your own name and QTH are looked up on QRZ.com — the one lookup the program never used to make. Both stay editable, and what you type is not overwritten by an answer arriving later.
- **Changed:** the net log window opens larger, and its two paragraphs of small print became one, to pay for the height the header takes. About eleven check-ins are visible at once.

0.16.0 — 22 August 2026

- **New:** a **net log** — a third window holding the stations that have checked in, numbered in the order they called, with the time, the call sign, and the operator's name and town. Stations go on it by ticking a box in the entry editor; unticking takes them off and the numbers close up.
- **New:** correcting a call sign corrects the check-in without renumbering the net.
- Save it before closing: this is the one record the program does not write to disk by itself, so that next week's net does not open with this week's still in it.

0.15.0 — 22 August 2026

- **Changed:** NoHands is now **NoHandsHam** — the name of the company that makes it. Same program, same people; the window, the Start Menu, the install folder, the manuals and the ADIF exports all say it.
- Installs over an existing NoHands where it stands: one entry in *Settings › Apps*, which changes its name. Your licence needs nothing re-entered.

0.14.1 — 22 August 2026

- **New:** correcting a call sign corrects **every** entry that named that station, along with the call sign log and the export. A station mis-heard once in a round was mis-heard the same way in the others, and correcting six entries by hand to reach a state the log had already assumed was six times the work.

0.14.0 — 21 August 2026

- **Changed:** the update-check address came off the settings pane. It was a box holding one value that only a release can change; what is left is the tick, which was the only real choice on that row.

0.13.3 — 20 August 2026

- **Fixed:** the update check asks this repository, which is public, rather than the source repository, which is not. Copies of 0.13.1 and 0.13.2 were asking a private repository and being told nothing at all.

0.13.2 — 20 August 2026

- No code changes — a version number and the documents stamped with it, so that 0.13.1 installs had a release to notice.

0.13.1 — 20 August 2026

- **Changed:** the update check asks what the newest *release* is, rather than reading a version file kept by hand. A file nothing in the build reads is a file that goes stale.

0.13.0 — 20 August 2026

- **Fixed:** the update check shipped in 0.12.3 could not reach the file it asks for, because that file was in a private repository. No code changed; the repository did.

0.12.3 — 20 August 2026

- **New:** the app says when a newer version has been published — one line at the foot of the settings pane, once a day, and nothing the rest of the time. Nothing is downloaded or installed for you.

0.12.2 — 20 August 2026

- **New:** the American police alphabet on request. US public-service nets spell with *Adam Boy Charles*, and those words used to mean nothing to the program — "William one Adam William" came out as four names with no call sign found.

0.12.1 — 20 August 2026

- **New:** the contact tally says how many are still to tick — `10/20 contacts in the log` — so the boxes waiting for attention are a number rather than a scroll.

0.12.0 — 19 August 2026

- **New:** both windows reopen where you left them, at the size you left them, and maximised if that is how you left them. A window whose monitor has gone comes back centred rather than somewhere unreachable.

0.11.0 — 19 August 2026

- **Changed:** the licence agreement is shown and has to be read before a licence seat is spent. Activating used to claim the seat the moment the button was pressed, with the agreement a PDF you opened only if you thought to.

0.10.0 — 19 August 2026

- **New:** a **Documents** pull-down in the window header reaching all five shipped documents — the guide, the manual, and the cookie, DMCA and installer-licence policies. They install alongside the program, so they open offline.

0.9.5 — 18 August 2026

- **Changed:** the app says it transcribes the radio, not the operator. The empty transcript window used to read "as you speak", which sent new users talking into a headset and getting a transcript of the wrong end of the QSO.
- **Fixed:** the README pointed at an `.msi` the release page no longer offers.

0.9.4 — 18 August 2026

- **Fixed:** every link in the window looks like a link. Three of the five — the OpenAI key page, the QRZ subscription and where to buy a licence — were still the dimmest text on screen.

0.9.3 — 18 August 2026

- **Fixed:** the two document links at the foot of the window were nearly unreadable.
- **Changed:** *Model* and *Base URL* moved down the settings pane, below the sliders an operator returns to mid-session. Both have a working default and were charging two rows of height to everything under them.

0.9.1 — 18 August 2026

- **New:** the guide and the manual are linked at the foot of the window. Both install with the program and get a Start Menu entry, and nothing in the window said so.

0.9.0 — 17–18 August 2026

- **New:** *Getting Started* — a short guide covering installing, getting a key and setting the four sliders, for people who should not have to read a 1,500-line manual first. Both PDFs now ship with every release.
- **New:** a failing session says *which* failure it is. A rejected API key, an expired licence and a dead microphone look identical from the operator's chair: Start is pressed and nothing arrives.
- **New:** every PDF is stamped with its version, so a manual on disk can be matched to the program that installed it.
- **Fixed:** the shipped binaries were four commits behind the release they were attached to.

0.8.5 — 16–17 August 2026

- **New:** the log export writes for **POTA, SOTA, Field Day** or as a plain generic log. It was POTA and nothing else, with POTA's own field hardcoded.
- **Changed:** the sliders that decide what is heard sit above the ones set once and left. The gate slider used to be several hundred pixels below the bottom of a laptop-sized settings pane.
- **Fixed:** text boxes and checkboxes got borders that can actually be found.

0.8.4 — 16 August 2026

- **Fixed:** field and checkbox edges were drawn at 1.3:1 against the window behind them — a boundary you have to find and click, rendered as a rumour. Now at the accessibility floor.

0.8.3 — 16 August 2026

- **Changed:** the level meter colours by where the bar is rather than by raw decibels. Nobody reads a meter in dBFS.

0.8.2 — 16 August 2026

- **New:** a live level meter under the gain slider. The slider shipped in 0.8.1 with nothing to set it against, and "turn the microphone up by 12 dB" is not something anybody knows in advance.
- **Fixed:** the desktop app opened a black command window beside itself on every launch.

0.8.1 — 16 August 2026

- **New:** a microphone gain slider, for an input that is too quiet before anything else in the app can help it — a headset across the room, or a receiver patched in on a cable.

0.8.0 — 16 August 2026

- **New:** the licence block says what the licence *is* — who it is issued to, how many computers, the expiry date, the trial days left — rather than one line summarising it and throwing the rest away.

0.7.2 — 15 August 2026

- **New:** the app says when transcription is failing, and checks the stored key at start-up rather than letting a session discover it. A session that transcribes nothing looks exactly like a dead microphone.

0.7.1 — 15 August 2026

- **Changed:** the licence moved to the top of *Accounts* and says whose it is. A bare "Licence" beside OpenAI and QRZ.com invited the guess that it belonged to one of them.

0.7.0 — 15 August 2026

- **New:** licensing. A key goes in the settings, *Activate* claims this computer's seat, and the app validates in the background at start-up — so a slow or unreachable licence server never delays the window.

0.6.7 — 15 August 2026

- **Changed:** the two account blocks fold away. Filled in on the first run and then left alone, they were costing a third of the settings pane for the life of the install.
- **Fixed:** the window asked for more height than the screen could show, which left the mouse and the picture out of step until it was resized by hand.

0.6.6 — 15 August 2026

- **New:** the settings say which accounts the program wants — OpenAI's, for the transcription you pay for, and QRZ.com's — and link to the page that issues each.

0.6.5 — 14 August 2026

- **New:** a running contact count at the foot of the transcript window. It used to appear only at the end of an export, after four dialogs.

0.6.4 — 14 August 2026

- **Fixed:** a window launched straight into Dark came up with its labels and help lines in the light theme's near-black — dim to the point of unreadable — until the look was changed and changed back.

0.6.3 — 14 August 2026

- **New:** a log opened from disk can be exported. Opening one used to put text in the window and nothing in the records, and the export refused.

0.6.2 — 14 August 2026

- **Changed:** the light/dark choice moved to the top right corner, out of the settings pane. It was the one setting you had to scroll to in order to use, and the only one whose effect is visible the instant it changes.

0.6.1 — 13 August 2026

- **New:** light or dark by choice, not only whatever Windows is set to. A laptop set light for daylight work is the wrong thing to have open in a dark shack.

0.6.0 — 13 August 2026

- **New:** a call sign you corrected or typed in is looked up on QRZ.com. It was the one call sign on screen that was certainly right and the only one with no name beside it, because the lookup that was made was made for the mis-hearing.

0.5.3 — 13 August 2026

- **New:** click an entry to correct what was heard, or **Add entry…** for a contact the microphone missed. Transcription is wrong in whole letters, and until now that letter was permanent.

0.5.2 — 13 August 2026

- **Fixed:** the version line added in 0.5.1 shipped invisible — drawn in the theme's *disabled* colour.

0.5.1 — 13 August 2026

- **New:** the version at the foot of the window. It was in four places, none of them open at the moment something goes wrong.

0.5.0 — 13 August 2026

- **Changed:** your own station came back off every transcript line. 0.3.0 put it there; it is the same twenty-odd characters on every row, and it belongs in the record rather than in the text.

0.4.0 — 12 August 2026

- **New:** export the ticked contacts as an **ADIF** file POTA will take. The app had written a real log since 0.3.0 and getting it uploaded still meant retyping every line.

0.3.0 — 12 August 2026

- **New:** your station is recorded with every contact — call sign, park or location, frequency and band, in four boxes across the top of the settings. A call sign says who was worked; nothing said where from.

0.2.2 — 12 August 2026

- **Changed:** the QRZ answer moved to the end of its entry rather than a line below it, where it cost a row per answered call sign and read as a second entry.

0.2.1 — 12 August 2026

- The version the program reports, the number stamped into the MSI and the file names in the manual moved together, from one constant.

0.2.0 — 12 August 2026

- **Changed:** HamLog is now **NoHands** — the module, both binaries, the installer and its product and folder names.

0.1.5 — 12 August 2026

- **Changed:** the installer ships as an `.exe` wrapping the `.msi`. An `.exe` is what people expect to download, and what a browser or mail client is least likely to hold up.

0.1.4 — 12 August 2026

- **Fixed:** the installer shortcut is in the install folder as well as the Start Menu, so an installed copy's own directory has a way back to it.

0.1.3 — 12 August 2026

- **New:** the QRZ answer appears in the transcript window, which is the window an operator actually watches. It was only ever reported on the control window.

0.1.2 — 12 August 2026

- **Changed:** the call sign log records only the contacts you have ticked, and names them from QRZ. A mis-heard run of phonetics can produce a perfectly well-formed call sign that no station holds.

0.1.1 — 12 August 2026

- **New:** a Windows MSI installer. The program was two loose executables and a PDF, and installing it meant knowing which one to run.

0.1.0 — 12 August 2026

- First release. A Windows desktop app and a CLI that transcribe continuously through the OpenAI speech API, built around what a radio operator actually needs: phonetics read back as call signs, a log of who was worked, and no typing.